
Politeness Is Not a Ruler: The Non-Linear Geometry of a Graded Social Trait in LLM Representations

Enrico Adamo
Bocconi University
enrico.adamo@studbocconi.it

Giovanni Berlinghieri
Bocconi University
giovanni.berlinghieri@studbocconi.it

Gabriele Bettineschi
Bocconi University
gabriele.bettineschi@studbocconi.it

Francesco Braicovich
Bocconi University
francesco.braicovich@studbocconi.it

Andrea Porta
Bocconi University
andrea.porta@studbocconi.it

Abstract

According to the linear representation hypothesis, LLMs encode high-level concepts as directions in activation space. The supporting evidence, however, comes almost entirely from *binary* contrasts, which can show a separating direction exists without testing whether intensity along it is linear. We ask this stronger question for *politeness* (a graded, high-level pragmatic trait with no external scalar) by adding an explicit NEUTRAL level between the poles, the minimal structure that makes linearity falsifiable. We build a dataset of same-content paraphrases varying only in politeness and probe Gemma 2 2B under several complementary geometries. **Politeness admits a clean binary axis yet is *not* a single ordered ruler: the levels are linearly orderable, but the adjacent transitions do not align and the neutral level lies off the line joining the poles—linear encoding of a graded trait need not be one-dimensional.**

1 Introduction

Linear representations in LLMs. A growing body of mechanistic work suggests that high-level concepts in large language models (LLMs) are encoded as approximately *linear* directions in the residual stream [Elhage et al., 2021, Park et al., 2024]. The idea predates LLMs (word-vector analogies [Mikolov et al., 2013] placed relations such as $\text{KING} - \text{MAN} + \text{WOMAN} \approx \text{QUEEN}$ on linear subspaces) and has since been sharpened for transformers, with linear codes reported for sentiment [Tigges et al., 2023], truthfulness [Marks and Tegmark, 2024], and continuous quantities such as space and time [Gurnee and Tegmark, 2024]. Representation-engineering and steering methods exploit the same structure to amplify or suppress behaviours without fine-tuning [Zou et al., 2023, Turner et al., 2023, Subramani et al., 2022, Rinsky et al., 2024].

Research Goal. Two clusters are collinear in any geometry, so a binary separation shows only that a direction *exists*; it cannot test whether intensity along it is linear—whether the representation behaves like an ordered *ruler*. We ask this stronger question for *politeness*: a high-level pragmatic attribute, graded by construction (a request can be rude, neutral, or deferential), socially constructed rather than anchored to an external scalar, and—unlike sentiment—only weakly carried by individual

valenced words. Concretely: *is politeness intensity encoded as a single, linearly graded direction, with the neutral level lying proportionally between the two poles?*

Contributions. Two commitments separate our study from prior work. *First, we test graded linearity, not binary separability* adding an explicit NEUTRAL level that makes linearity falsifiable. *Second, we probe a high-level pragmatic trait, not a lexical or externally grounded one.* Both demand strict *content control* so we build the dataset ourselves (Section 3.1) with an independent LLM judge.

2 Background and Related Work

Measuring linearity, not just finding directions. The works above show that concept directions *exist*; our question is about their *geometry*, which raises issues those studies do not face: a linear probe recovering an attribute is weak evidence that the model actually *uses* that direction [Belinkov, 2022, Alain and Bengio, 2016], so difference-of-means estimators and causal steering are preferred for isolating behaviourally relevant axes [Marks and Tegmark, 2024, Rinsky et al., 2024]. Measuring whether intensity is encoded *linearly* is subtler still: cosine similarity in the raw residual stream is not coordinate-free [Park et al., 2024] and the space is strongly anisotropic [Ethayarajh, 2019], so we evaluate directions under several complementary geometries rather than one (Section 3).

Politeness in NLP: behavioural, not representational. Politeness has been studied extensively in NLP, but almost exclusively at the *behavioural* level—phrase-level classifiers [Danescu-Niculescu-Mizil et al., 2013], controllable rewriters [Niu and Bansal, 2018, Madaan et al., 2020], and few-shot style transfer [Reif et al., 2022]—all operating on mode outputs, not internal representations.

Graded representations, and the open question. Tigges et al. [2023] and Konen et al. [2024] recover single directions for sentiment and emotion, but both treat them as bipolar and neither tests whether an intermediate level sits proportionally between the poles. Heinzerling and Inui [2024] come closest—distinguishing *linear* from *monotonic* encodings—but only for externally grounded numeric properties that vary across *entities*, whereas politeness has no external scalar and must be elicited by rewriting the *same* content. Whether a graded *pragmatic* trait forms a one-dimensional ordered ladder thus remains open.

3 Method

3.1 Dataset construction

Our stimulus set must be *content controlled*: items at different politeness levels must differ in intensity and, as far as possible, nothing else. No public corpus offers this at the granularity we need, so we generate the benchmark ourselves with a two-model pipeline: a *generator* (Gemini 2.0 Flash) and an independent *judge* (DeepSeek v4 Flash). Appendix A gives full detail.

Scenarios and ladders. The unit of generation is a *scenario*: a fixed communicative situation that pins down everything except how politely it is expressed—an *intent* (one of ten illocutionary types: request, refusal, disagreement, criticism, bad-news delivery, apology, complaint, reminder, sensitive inquiry, correction), a domain, the speaker–addressee relation, and an *intent target* (the concrete content every rung must convey, e.g. “send the budget spreadsheet by end of day”). From each scenario the generator writes a *ladder*: three sentences, one per signed level (*negative*, *neutral*, *positive*, neutral a genuine zero), conveying the same target and differing only in politeness, with up to three paraphrases per level. Ten intents keep any single act from dominating the geometry; the dataset is English-only, with 72 complete scenarios (633 paraphrases) after filtering.

Confound control and filtering. Because politeness is read from the *difference* between a scenario’s rungs, content varying across rungs would be mistaken for trait signal. We guard against this twice over. By construction, paraphrases within a level use distinct *cue families* (courtesy markers, indirectness, gratitude, imposition acknowledgment, impersonal phrasing), and every rung stays within a fixed word-count tolerance, so neither a single marker nor length can cue the level. The judge then applies a *content* check (faithfulness to the target), a *length* check, and—decisively—a *blind*

Estimator	Bal. acc. raw / centered	Cosine to mean-diff	Descriptor	Observed (95% CI)	p
Difference-of-means	0.914 / 0.971	1.00	Step cosine	-0.23 [-0.33, -0.13]	0.001
k -means	0.759 / 0.967	0.996	Apex angle	77° [71, 83]	—
Logistic regression	0.962 / 0.991	0.787	Midpoint residual	0.61 [0.56, 0.68]	0.001
PCA	0.755 / 0.950	0.966			
Random direction	—	0.025			

Table 1: Left: binary POSITIVE-vs-NEGATIVE contrast at the last token of layer 13, reporting held-out balanced accuracy under scenario-grouped CV and cosine alignment to the centered difference-of-means axis. Right: triple geometry at layer 13 over 72 scenarios, with scenario-bootstrap CIs and noise-null p values. The binary axes separate the poles after centering, while the triple descriptors remain far from a straight ladder geometry.

intensity check that re-rates politeness with level labels hidden, keeping a bundle only if its blind per-level means are ordered NEGATIVE < NEUTRAL < POSITIVE with a minimum gap. The ordering is thus recovered by a label-blind rater, not asserted. A bag-of-words baseline already classifies the level at 85.0% (chance 33%)—a few lexical cues carry much of the signal—so we do not claim separability as a result and study the *graded geometry* instead, which it cannot express.

4 Experiments

We probe **Gemma 2 2B** [Gemma Team, 2024], reading the residual stream at layer 13 of 26—mid-network, where high-level features tend to concentrate [Tigges et al., 2023]—at the *last token* of each paraphrase, the convention in steering work [Zou et al., 2023, Tigges et al., 2023]; mean-pooling reproduces every result (Section C). One methodological choice matters throughout: residual activations are dominated by each scenario’s content, so we treat `scenario_id` as a block, subtracting its across-level mean and evaluating held-out under GroupKFold. This *blocked* estimator is the one we interpret; a naive pooled baseline (Section B) shows how much content otherwise inflates the numbers. We re-measure every alignment under several whitenings and against a within-scenario permutation null, and report nothing the noise floor can explain (Section B). Because a bag-of-words classifier already labels the level at 85.0% (chance 33%), the question is never whether politeness is *decodable*, but whether it is laid out as a single linear scale. We ask this in three steps of increasing strength: does a binary axis exist (Section 4.1); do the three ordered levels fall on one straight line (Section 4.2); and does a single direction steer politeness across every context (Section 4.3)? All analyses use the 72 complete scenarios (633 paraphrases).

4.1 A clean binary axis exists

Restricting the ladder to its poles reproduces, for politeness, the linear separability Tigges et al. [2023] reported for sentiment. The two poles project onto the difference-of-means axis with ROC-AUC 0.98 and Cohen’s $d = 3.0$, and a logistic probe classifies held-out scenarios at 96.2% balanced accuracy (97.1% after within-scenario centering)—above the GPT-2 sentiment reference (0.80–0.89) and far above a random-direction floor (cosine ≈ 0.02). Four independent estimators (difference-of-means, k -means, logistic, PCA) recover essentially the same axis once content is controlled (Table 1). Both adjacent half-steps are separable on their own too (NEG-NEU 88%/95% raw/centered; NEU-POS 76%/85%), so each transition is individually linearly recoverable. Binary separability is therefore not in question—which is exactly why it cannot, by itself, establish a *graded ruler*.

4.2 The graded ladder is ordered but bent

Adding NEUTRAL lets us ask whether the three levels lie, *in order*, along one line. They are reliably *ordered*: against the endpoint (negative→positive) axis the blocked Spearman $\rho = 0.84$ and Kendall $\tau = 0.70$ (near their structural ceilings of 0.94/0.83), a cross-validated ridge decodes rank at $R^2 = 0.85$, and 81% of scenarios are monotone—every metric beats its within-scenario permutation null at $p \approx 0.01$. But the levels are *not collinear*. The two transition directions (NEU-NEG and POS-NEU) point apart rather than together: their cosine is -0.23 (a straight ladder predicts +1), the apex angle at neutral is 77° (straight = 180°), and the neutral centroid sits a normalised 0.61 of

the pole-to-pole span off the midpoint (straight = 0). The bend survives a stringent null: forcing neutral onto the true midpoint and re-injecting the empirical paraphrase noise yields a null step-cosine of $+0.92 \pm 0.01$ and midpoint residual 0.10 ± 0.01 , so measurement noise cannot manufacture the observed values ($p = 0.001$, $\approx$ tens of σ ; ??). The negative cosine persists under all four inner products (-0.14 to -0.23) and after disattenuation. The ladder is thus a *V-shape*: orderable, yet not laid out along a single straight ruler.

4.3 No universal steering direction

Steering presumes a *single* direction that moves the trait everywhere; two findings rule that out. Within a scenario the ladder is bent (Section 4.2): because $\text{neg} \rightarrow \text{neu}$ and $\text{neu} \rightarrow \text{pos}$ point in substantially different directions (77° apex), no single axis traverses the whole scale and the $\text{neg} \rightarrow \text{pos}$ vector skips the neutral level rather than passing through it. *Across* scenarios the contrast is not shared either: a global politeness axis is highly reproducible on average (split-half Spearman–Brown $R = 0.94$), yet recovers only about a fifth of any individual scenario’s $\text{neg} \rightarrow \text{pos}$ contrast (mean capture $\cos^2 = 0.20$, IQR $0.13\text{--}0.26$). The shortfall is not estimation noise. Part of it is the bend itself: the off-line direction along which neutral is displaced from the pole-to-pole line is reliable and shared (neutral falls on the same side in 96% of scenarios, $R = 0.89$), yet spanning the full plane of the politeness axis and this bend direction still captures only 0.21 of within-scenario variance. The rest is genuine, scenario-specific structure.

5 Conclusion

We asked whether a graded pragmatic trait—politeness—is encoded as a single linear ruler in Gemma 2 2B. Using a content-controlled, blind-validated ladder with an explicit neutral level, we separate three senses of “linear representation” that binary studies conflate. *Ordinal* linearity **holds**: the levels sit in order along a recoverable axis when projected on it. *Metric* linearity **fails**: neutral lies off the pole-to-pole line and adjacent steps anti-align at a 77° apex, so the scale is not one straight ruler. *Universal-direction* linearity **fails**: no single vector reproduces the contrast across contexts. Linear decodability of a graded trait thus does not entail a one-dimensional, steerable encoding—a distinction with direct consequences for activation steering, which a single politeness vector moves only $\sim 20\%$ of the way in any given context.

References

- Guillaume Alain and Yoshua Bengio. Understanding intermediate layers using linear classifier probes. *arXiv preprint arXiv:1610.01644*, 2016.
- Yonatan Belinkov. Probing classifiers: Promises, shortcomings, and advances. *Computational Linguistics*, 48(1):207–219, 2022.
- Cristian Danescu-Niculescu-Mizil, Moritz Sudhof, Dan Jurafsky, Jure Leskovec, and Christopher Potts. A computational approach to politeness with application to social factors. In *Proceedings of the 51st Annual Meeting of the Association for Computational Linguistics (ACL)*, pages 250–259, 2013.
- Nelson Elhage, Neel Nanda, Catherine Olsson, Tom Henighan, Nicholas Joseph, Ben Mann, Amanda Askell, Yuntao Bai, Anna Chen, Tom Conerly, et al. A mathematical framework for transformer circuits. In *Anthropic Transformer Circuits Thread*, 2021.
- Kawin Ethayarajh. How contextual are contextualized word representations? comparing the geometry of BERT, ELMo, and GPT-2 embeddings. In *Proceedings of the 2019 Conference on Empirical Methods in Natural Language Processing and the 9th International Joint Conference on Natural Language Processing (EMNLP-IJCNLP)*, pages 55–65, 2019.
- Gemma Team. Gemma 2: Improving open language models at a practical size. *arXiv preprint arXiv:2408.00118*, 2024.
- Wes Gurnee and Max Tegmark. Language models represent space and time. In *International Conference on Learning Representations (ICLR)*, 2024.

- Benjamin Heinzerling and Kentaro Inui. Monotonic representation of numeric properties in language models. In *Proceedings of the 62nd Annual Meeting of the Association for Computational Linguistics (Short Papers)*, 2024.
- Kai Konen, Sophie Jentzsch, Diaoulé Diallo, Peer Schütt, Oliver Bensch, Roxanne El Baff, Dominik Opitz, and Tobias Hecking. Style vectors for steering generative large language models. In *Findings of the Association for Computational Linguistics: EACL 2024*, pages 782–802, 2024.
- Aman Madaan, Amrith Setlur, Tanmay Parekh, Barnabas Poczos, Graham Neubig, Yiming Yang, Ruslan Salakhutdinov, Alan W. Black, and Shrimai Prabhumoye. Politeness transfer: A tag and generate approach. In *Proceedings of the 58th Annual Meeting of the Association for Computational Linguistics (ACL)*, pages 1869–1881, 2020.
- Samuel Marks and Max Tegmark. The geometry of truth: Emergent linear structure in large language model representations of true/false datasets. In *Conference on Language Modeling (COLM)*, 2024.
- Tomas Mikolov, Ilya Sutskever, Kai Chen, Greg S. Corrado, and Jeff Dean. Distributed representations of words and phrases and their compositionality. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2013.
- Tong Niu and Mohit Bansal. Polite dialogue generation without parallel data. In *Transactions of the Association for Computational Linguistics (TACL)*, volume 6, pages 373–389, 2018.
- Kiho Park, Yo Joong Choe, and Victor Veitch. The linear representation hypothesis and the geometry of large language models. In *Proceedings of the 41st International Conference on Machine Learning (ICML)*, 2024.
- Emily Reif, Daphne Ippolito, Ann Yuan, Andy Coenen, Chris Callison-Burch, and Jason Wei. A recipe for arbitrary text style transfer with large language models. In *Proceedings of the 60th Annual Meeting of the Association for Computational Linguistics (ACL), Short Papers*, pages 837–848, 2022.
- Nina Rimskey, Nick Gabrieli, Julian Schulz, Meg Tong, Evan Hubinger, and Alexander Turner. Steering Llama 2 via contrastive activation addition. In *Proceedings of the 62nd Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 15504–15522, 2024.
- Nishant Subramani, Nivedita Suresh, and Matthew E. Peters. Extracting latent steering vectors from pretrained language models. In *Findings of the Association for Computational Linguistics: ACL 2022*, pages 566–581, 2022.
- Curt Tigges, Oskar J. Hollinsworth, Atticus Geiger, and Neel Nanda. Linear representations of sentiment in large language models. In *arXiv preprint arXiv:2310.15154*, 2023.
- Alexander Matt Turner, Lisa Thiergart, Gavin Leech, David Udell, Juan J. Vazquez, Ulisse Mini, and Monte MacDiarmid. Activation addition: Steering language models without optimization. In *arXiv preprint arXiv:2308.10248*, 2023.
- Andy Zou, Long Phan, Sarah Chen, James Campbell, Phillip Guo, Richard Ren, Alexander Pan, Xuwang Yin, Mantas Mazeika, Ann-Kathrin Dombrowski, Shashwat Goel, Nathaniel Li, Michael J. Byun, Zifan Wang, Alex Mallen, Steven Basart, Sanmi Koyejo, Dawn Song, Matt Fredrikson, J. Zico Kolter, and Dan Hendrycks. Representation engineering: A top-down approach to AI transparency. *arXiv preprint arXiv:2310.01405*, 2023.

A Dataset construction details

This appendix gives the full specification of the corpus summarised in Section 3.1. The pipeline runs in four stages—scenario generation, paraphrase generation, judging/filtering, and activation extraction—driven by two models, a *generator* (Gemini 2.0 Flash) and a *judge* (DeepSeek v4 Flash).

A.1 Construct rubric

Politeness is defined on a signed three-point scale as the mitigation of face threat, deference, and social consideration, with neutral a true zero (no courtesy and no rudeness markers) and the poles equal-and-opposite departures from it. **Negative** is dismissive, curt, or impatient, with rudeness from tone and framing rather than profanity, threats, or personal attacks (which would be lexical giveaways); **neutral** is plain and even-toned; **positive** is respectful and mitigated—appreciation, deference, acknowledgment of the imposition—through framing rather than added words. Held constant across the three levels of a scenario: the intent target, named entities, deadlines, the polarity of the act, the amount imposed, and the length.

A.2 Scenario schema and intents

Each scenario is a structured record specifying the *intent* (one of ten: *request*, *refusal*, *disagreement*, *criticism_or_feedback*, *bad_news_delivery*, *apology*, *complaint*, *reminder*, *inquiry_sensitive*, *correction*), a *domain* (e.g. workplace, restaurant, healthcare), the *audience relation* between speaker and addressee (e.g. peer, subordinate, service agent), a *communicative goal*, the *intent target* (the concrete content every rung must convey), and a *target word count*. Each intent carries its own invariance constraint—for example, a refusal must remain a refusal and may not become a partial acceptance. Intents are drawn from a shuffled queue, so the generated set is balanced across the ten.

A.3 Generation and judging

Models. Two models drive the pipeline (Table 2): the generator produces both scenarios and paraphrases, and the judge performs all filtering, including the blind intensity rating. A configuration check refuses to run if generator and judge share a model family, so generation and validation are never performed by the same model.

Role	Model	Temperature
Generator (scenarios + paraphrases)	google/gemini-2.0-flash-001	0.5
Judge (filtering + blind scoring)	deepseek/deepseek-v4-flash	0.0
Probed model (activation extraction)	google/gemma-2-2b	—

Table 2: Models used in dataset construction. Generator and judge are required to be from different families; the probed model is independent of both.

Acceptance criteria. After generation, each scenario’s bundle of paraphrases passes through three judge checks, with parameters fixed for the run:

- **Context (content preservation).** The judge scores each paraphrase in $[0, 1]$ for how faithfully it conveys the scenario’s intent target—explicitly ignoring politeness and fluency—and paraphrases must meet a minimum acceptance score of 0.70.
- **Trait intensity (blind monotonicity).** The judge re-rates each paraphrase’s politeness in $[0, 1]$ with level labels *hidden*: texts appear under opaque codes (q000, q001, ...) in shuffled order, with no indication of which belong together. A bundle is kept only if the blind per-level means are ordered $\text{NEGATIVE} < \text{NEUTRAL} < \text{POSITIVE}$ with a minimum adjacent gap of 0.10.
- **Length.** Per-level word counts must not diverge by more than a ratio of 1.15, so length cannot cue the level.

A near-duplicate guard additionally removes paraphrases whose content overlaps an earlier one above a Jaccard threshold. The intent queue and the blind-scoring order use a fixed random seed for reproducibility.

A.4 Resulting corpus

From 100 generated scenarios (900 paraphrases), 633 passed all three checks (a 70.3% acceptance rate); the 267 rejections were dominated by the trait-intensity gate (163), followed by length (72)

and context (32). The filtered corpus spans 72 scenarios for which all three levels survive, balanced across the ten intents, with per-level counts of 212 negative, 214 neutral, and 207 positive paraphrases (Table 3). Sentence length is well matched across levels (mean word counts 16.9, 17.4, 17.7), confirming the length control held in the final data. Cue families are recorded as free-text labels and span a broad vocabulary rather than a fixed set, reflecting the diversification target.

Quantity	Value
Scenarios generated	100
Paraphrases generated	900
Paraphrases passing all checks	633 (70.3%)
rejected: trait intensity / length / context	163 / 72 / 32
Complete scenarios (all three levels) in filtered set	72
Paraphrases by level (neg / neu / pos)	212 / 214 / 207
Mean length by level (neg / neu / pos, words)	16.9 / 17.4 / 17.7
Bag-of-words classifier accuracy (5-fold CV)	85.0% (chance 33%)

Table 3: Dataset statistics for politeness.

Lexical baseline. The bag-of-words baseline uses unigram and bigram counts (min document frequency 2) with logistic regression under stratified 5-fold cross-validation, predicting the level from raw text. It reaches 85.0% accuracy (chance 33%, $n = 633$): politeness is strongly lexically cued, so we report this not as a rejection criterion but as a floor the geometric analysis must exceed to claim structure beyond surface keywords. The most discriminative n -grams per level are intuitive—NEGATIVE: *immediately, no, late, wrong, seriously, not*; NEUTRAL: *please, scheduled, soon, should*; POSITIVE: *appreciate, hope, thanks, perhaps, might*—confirming that the surface signal is exactly the courtesy and curtness vocabulary one would expect.

A.5 Activation extraction

Activations are taken from the probed model’s *prefill* pass over each paraphrase—no text is generated—at every layer in 1–22, and cached. Each paraphrase is reduced to one vector per layer by token pooling: the *last* content token (our primary setting) or the *average* over content tokens (excluding the beginning-of-sequence, end-of-sequence, and padding tokens), with both poolings stored. Paraphrases sharing a (trait, level, scenario) key are averaged. The pooling-invariant unembedding covariance, used for the causal-geometry analysis, is computed once over the model’s full output-embedding matrix and stored alongside the activations.

A.6 Example bundle

Field	Value
<code>scenario_id</code>	politeness-request-001
Intent / domain	request / workplace
Intent target	send the latest budget spreadsheet by end of day
NEGATIVE	“Is it too much to ask for the budget spreadsheet by day’s end? Get it done.”
NEUTRAL	“Please provide the latest budget spreadsheet before the end of the day; it’s needed today.”
POSITIVE	“Would you mind sending over the latest budget spreadsheet by the end of the day, please?”

Table 4: One scenario (politeness-request-001) realised at all three levels. All three rungs convey the same intent target—*send the latest budget spreadsheet by end of day*—and differ only in politeness.

B Estimation, geometry, and significance

Blocked vs. naive-pooled estimators. Residual activations carry large variance unrelated to politeness, dominated by the invariant content of each scenario. Since `scenario_id` fixes that content, we

treat the scenario as a block and report every metric two ways. The *blocked* (fixed-effects) estimator subtracts each scenario’s across-level mean before any geometry is measured and evaluates held-out under GroupKFold, so a scenario’s levels never split across train and test; it is the estimator we interpret. The *naive pooled* estimator pools points by level only and serves as a baseline that exposes how much the content confound inflates the numbers—blocking raises the Spearman ρ from 0.77 to 0.84 and the probe R^2 from 0.77 to 0.85, and the midpoint residual moves from a pooled 0.61 to a 0.72 per-scenario median.

Four inner products and disattenuation. A near-zero or near-one raw cosine can be an artefact of coordinates rather than a fact about politeness, so we re-measure the transition-direction alignment under four inner products: raw Euclidean; diagonal (anisotropy) whitening; the causal whitening of Park et al. [2024], from the inverse square root of the unembedding covariance over the full vocabulary; and Mahalanobis whitening from the Ledoit–Wolf-shrunk pooled within-level covariance. The adjacent-step cosine stays negative in all four (−0.23 Euclidean, −0.20 anisotropy, −0.21 causal, −0.14 Mahalanobis), and disattenuating by split-half reliability (300 random halves, rescaled by Spearman–Brown) leaves it negative (−0.25 to −0.14). The bend is thus neither a coordinate nor an estimation-noise artefact.

Within-scenario permutation null. To test whether the geometry could arise by chance we shuffle the three level labels *inside* each scenario, holding the block structure intact, and recompute every metric over N permutations; a global shuffle would also destroy content and topic structure, a much weaker null. We report add-one Monte-Carlo p -values, $p = (1 + \#\{\text{null beyond observed}\})/(1 + N)$, with the appropriate tail per metric (lower for the midpoint residual, upper otherwise). All six ordinal metrics clear the null at $p \approx 0.01$.

Lexical baseline and layer robustness. A bag-of-words classifier (unigram and bigram counts, min document frequency 2, logistic regression, stratified 5-fold CV) predicts the level at 85.0% (chance 33%): politeness is strongly lexically cued, so we treat this as a floor the geometry must exceed, not a result. The most discriminative n -grams are the expected courtesy/curtness vocabulary (NEG: *now, immediately, no, late*; NEU: *please, scheduled, should*; POS: *appreciate, perhaps, hope, thanks*). Repeating the direction-agreement, alignment, and full-distribution analyses across layers 1–22 gives the same qualitative picture throughout: the ordinal structure and the off-line bend are properties of the model, not of the fixed layer.

C Full results under both token poolings

The main text reads the last-token residual. As a robustness check we re-run *every* analysis—binary, graded, and steering—under *mean-pooling*, averaging the layer-13 residual over each paraphrase’s content tokens (excluding the BOS, EOS, and padding positions). The mean-pooled bundle is precomputed under `data/<timestamp>/representations/avg_token/` and consumed by the same scripts (a `TOKEN_POOLING` constant selects the branch); figures land under `results/<timestamp>/.../avg_token/`. This appendix reports the complete numerical and graphical results for *both* poolings side by side, organised by the three claims of the main text. Every table carries last-token and mean-pooled columns; every figure family is shown for both poolings.

The headline is that every conclusion holds under mean-pooling, and is if anything sharper than at the last token. The binary poles separate at least as well (logistic balanced accuracy 0.976 raw, ≥ 0.99 after centering). The ladder is ordered (blocked Spearman $\rho = 0.90$, Kendall $\tau = 0.76$, probe $R^2 = 0.92$) yet bent: step cosine −0.29, apex 73° , midpoint residual 0.66, with the same noise null rejected at $p = 0.001$. And there is still no universal steering direction: the mean axis is reproducible ($R = 0.96$) but captures only $\cos^2 = 0.25$ of any scenario’s contrast, the shared plane captures 0.26 (94% of the two-dimensional ceiling), and the tilt is intent-structured (within 0.39 vs. between 0.22). Reading a single token is therefore not what produces the non-linearity.

C.1 Binary politeness axis

Table 5 gives held-out balanced accuracy for all four estimators on all three pairwise contrasts, raw and after within-scenario centering, for both poolings. Centering—which removes the scenario-content confound—lifts every estimator to near-ceiling, confirming that a clean binary axis exists at each

Contrast	Method	Last token		Mean pooling	
		raw	centered	raw	centered
pos vs. neg	MeanDiff	0.914	0.971	0.955	0.995
	KMeans	0.759	0.967	0.495	0.995
	LogReg	0.962	0.991	0.976	0.998
	PCA	0.755	0.950	0.679	0.993
neu vs. neg	MeanDiff	0.881	0.953	0.904	0.977
	KMeans	0.763	0.951	0.686	0.974
	LogReg	0.929	0.967	0.963	0.995
	PCA	0.763	0.941	0.732	0.969
pos vs. neu	MeanDiff	0.757	0.848	0.787	0.920
	KMeans	0.520	0.802	0.507	0.901
	LogReg	0.843	0.950	0.879	0.960
	PCA	0.529	0.788	0.504	0.898

Table 5: Binary balanced accuracy (5-fold, layer 13) for every estimator, contrast, and pooling, raw and after within-scenario centering. Chance is 0.50; Tigges et al. report .78–.89 for the pos-vs-neg contrast.




Figure 1: Binary balanced accuracy across layers. *Left*: last token. *Right*: mean pooling. The axis is recoverable throughout the network, with a broad plateau around the focal layer 13.

adjacent step as well as across the full range. The hardest contrast is the positive-vs-neutral half-step, exactly the half-step that the graded analysis later finds bent. Figure 1 shows the accuracy is stable across layers, and Figures 2 to 5 give the per-contrast direction-agreement and projection diagnostics.




Figure 2: Direction agreement (last token), columns neu-vs-neg, pos-vs-neu, pos-vs-neg. *Top:* raw. *Bottom:* within-scenario centered.




Figure 3: Direction agreement (mean pooling), same layout as Figure 2.




Figure 4: Projection onto the difference-of-means axis (last token), columns neu-vs-neg, pos-vs-neu, pos-vs-neg. *Top*: raw. *Bottom*: centered.




Figure 5: Projection onto the difference-of-means axis (mean pooling), same layout as Figure 4.

Descriptor	Last token point [95% CI]	Mean pooling point [95% CI]
Step cosine	-0.229 [-0.331, -0.127]	-0.289 [-0.376, -0.212]
Apex angle (°)	76.8 [70.7, 82.7]	73.2 [67.9, 77.8]
Midpoint residual	0.613 [0.560, 0.679]	0.661 [0.612, 0.723]
Markedness	0.537 [0.487, 0.617]	0.619 [0.571, 0.686]
Markedness ratio	1.073 [0.974, 1.233]	1.238 [1.141, 1.372]
Valence offset	0.297 [0.214, 0.346]	0.232 [0.159, 0.294]
Effective vertices	4.69 [4.35, 5.09]	4.92 [4.63, 5.30]

Table 6: Pooled triple-geometry descriptors with scenario-bootstrap intervals (1000 resamples), layer 13. A straight ladder would give step cosine +1, apex 180°, and midpoint residual 0.

Metric	observed		null mean (sd)		<i>p</i>
	last	avg	last	avg	
Step cosine	-0.229	-0.289	0.919 (0.015)	0.946 (0.010)	0.001
Midpoint residual	0.613	0.661	0.104 (0.010)	0.084 (0.008)	0.001

Table 7: Linear-ladder noise null ($N = 1000$): Gaussian noise calibrated to the observed scatter, placed on a perfectly straight ladder, cannot reproduce either the negative step cosine or the large midpoint residual, under either pooling.

C.2 Graded-ladder geometry

Table 6 reports the seven geometry descriptors with scenario-bootstrap 95% intervals for both poolings, and Table 7 the linear-ladder noise null. The bend is unambiguous in both: the adjacent steps anti-align (step cosine $-0.23/-0.29$), the apex sits near a right angle rather than the straight-ladder 180°, and the midpoint residual is far above what noise on a straight ladder can produce ($p = 0.001$). Table 8 breaks the geometry down by communicative intent: every one of the ten intents is bent (apex 59–90°, none near 180°), with requests and disagreements the straightest and complaints, corrections, and refusals the most bent under both poolings. Figures 6 and 7 show the bootstrap and noise-null distributions; Figure 8 shows the per-intent triangles.

Intent	n	Last token			Mean pooling		
		Apex ^o	Step cos	Mid.	Apex ^o	Step cos	Mid.
apology	8	78.2	-0.20	0.60	84.0	-0.10	0.55
bad-news-delivery	4	72.3	-0.30	0.64	59.3	-0.51	0.83
complaint	7	66.2	-0.40	0.76	63.7	-0.44	0.80
correction	9	67.9	-0.38	0.73	74.5	-0.27	0.66
criticism-or-feedback	8	67.9	-0.38	0.73	70.9	-0.33	0.69
disagreement	6	82.9	-0.12	0.56	90.3	+0.01	0.50
inquiry-sensitive	4	79.8	-0.18	0.60	74.6	-0.27	0.66
refusal	8	73.7	-0.28	0.65	63.4	-0.45	0.80
reminder	6	73.2	-0.29	0.66	71.2	-0.32	0.68
request	12	88.6	-0.02	0.51	82.7	-0.13	0.57
ALL	72	76.8	-0.23	0.61	73.2	-0.29	0.66

Table 8: Triple geometry pooled within each communicative intent, both poolings (layer 13). Apex angles stay well below the straight-ladder 180° for every intent.




Figure 6: Triple geometry, last token. *Left*: bootstrap distributions of the geometry descriptors against the straight-ladder and pentagon references. *Right*: the linear-ladder noise null.




Figure 7: Triple geometry, mean pooling; same panels as Figure 6.




Figure 8: Per-intent triangles. *Left*: last token. *Right*: mean pooling. Every intent's triple is bent at an interior apex.

Metric	Last token		Mean pooling	
	pooled	within	pooled	within
Spearman ρ	0.775	0.841	0.812	0.895
Kendall τ	0.631	0.701	0.668	0.764
Probe R^2	0.769	0.853	0.791	0.916
Midpoint residual	0.613	0.717	0.661	0.727
PC1 fraction	0.782	0.782	0.715	0.715

Table 9: Ordinal-linearity metrics under the naive-pooled and blocked (within-scenario) estimators. Ordering is high and rises under blocking; the midpoint residual stays high, marking the bend.

Metric	observed		null mean		p	
	last	avg	last	avg	last	avg
Spearman ρ	0.841	0.895	-0.006	-0.006	0.010	0.010
Kendall τ	0.701	0.764	-0.005	-0.004	0.010	0.010
Probe R^2	0.853	0.916	-0.388	-0.409	0.010	0.010
Monotone fraction	0.806	0.889	0.345	0.310	0.010	0.010
Midpoint resid. median	0.717	0.727	0.844	0.852	0.010	0.010
PC1 fraction	0.782	0.715	0.600	0.619	0.010	0.050

Table 10: Within-scenario label-permutation null ($N = 100$, add-one p): level labels are shuffled inside each scenario, holding block structure intact. Every ordinal metric clears the null.

C.3 Ordinal structure and the bend

Table 9 reports the five ordinal-linearity metrics under the blocked (within-scenario) and naive-pooled estimators for both poolings; blocking raises every metric, and the within-scenario ladder is strongly ordered ($\rho = 0.84/0.90$, probe $R^2 = 0.85/0.92$) yet its midpoint residual stays high. Table 10 gives the within-scenario permutation null: all six metrics clear it at $p \approx 0.01$ (the sole exception is the mean-pooled PC1-fraction at $p = 0.05$, since the centroids are already nearly planar). Table 11 re-measures the adjacent-step cosine under five inner products, raw and disattenuated by split-half reliability; it remains negative in every space under both poolings, so the bend is neither a coordinate nor a reliability artefact. Table 12 lists the full step-vector cosine matrix. The supporting figures are Figures 9 to 14.

Inner product	Last token		Mean pooling	
	raw	disatt.	raw	disatt.
Euclidean (raw)	−0.229	−0.253	−0.289	−0.307
Anisotropy-corrected	−0.203	−0.229	−0.133	−0.146
Park causal	−0.208	−0.231	−0.278	−0.296
Mahalanobis / LDA	−0.135	−0.138	−0.192	−0.196
Within-subjects noise	−0.088	−0.089	−0.151	−0.152

Table 11: Adjacent-step cosine under five inner products, raw and disattenuated by split-half reliability (300 halves, Spearman–Brown). The bend survives every whitening and the reliability correction.

Step pair	Last token	Mean pooling
(pos−neu) · (neu−neg)	−0.229	−0.289
(pos−neu) · (pos−neg)	+0.354	+0.397
(neu−neg) · (pos−neg)	+0.829	+0.764

Table 12: Cosines among the three step vectors. The two adjacent half-steps anti-align, while each aligns with the full neg→pos contrast—the signature of a bent ladder.




Figure 9: Last token. *Left*: cosines among the three step vectors; the adjacent half-steps anti-align (−0.23). *Right*: the within-scenario permutation null for the six ordinal metrics.




Figure 10: Mean pooling; same panels as Figure 9.




Figure 11: Inner-product geometry. *Top*: last token. *Bottom*: mean pooling. *Left*: adjacent-step cosine across the five inner products with disattenuation. *Right*: the full step-pair cosine grid per inner product.




Figure 12: Last token. *Left*: PCA of the level centroids, raw vs. within-scenario centered. *Right*: a grid of per-scenario triples showing the recurring bend.




Figure 13: Mean pooling; same panels as Figure 12.




Figure 14: Layer robustness. *Top*: last token. *Bottom*: mean pooling. *Left*: ordinal-linearity metrics across layers. *Right*: step-vector cosines across layers. The ordering-with-a-bend pattern holds throughout.

Quantity	Last token	Mean pooling
Valence axis reliability R (full)	0.938	0.956
Mean $\cos(\text{contrast, global valence})$	0.428	0.493
Steering capture $\cos^2$ (layer 13)	0.196	0.252
capture IQR	[0.127, 0.258]	[0.186, 0.325]
Mean $\cos(\text{bend, markedness axis})$	0.327	0.413
Median principal angle θ_1 ($^\circ$)	59.3	55.3
Median principal angle θ_2 ($^\circ$)	73.3	67.7
Median plane overlap	0.175	0.235
Contrast $\cos$ within-intent	0.252	0.389
Contrast $\cos$ between-intent	0.164	0.216
within—between gap	0.088	0.174
Shared plane: valence-1D captured	0.139	0.167
Shared plane: 2D captured	0.206	0.264
Neutral markedness sign consistency	95.8%	98.6%
Shared-bend R^2 (CV)	0.112	0.167
Scenarios improved by shared bend	83.3%	84.7%
Participation ratio	26.5	20.0

Table 13: Steering and subspace decomposition, layer 13. A reproducible global axis coexists with low per-scenario capture, an intent-structured residual tilt, and a high-dimensional spread of per-scenario planes—no single universal steering direction.




Figure 15: No universal steering direction (last token). *Left*: per-scenario alignment to the global axis and plane-tilt angles. *Middle*: steering capture ($\cos^2$) across layers, plateauing near 0.20. *Right*: the shared (valence, markedness) plane.

C.4 Steering geometry and subspace dimensionality

Table 13 collects the steering decomposition for both poolings. Throughout, the *valence axis* is the global pole-to-pole politeness direction (negative→positive), and the *markedness axis* is the shared off-line direction along which the neutral level is displaced from that line. The global valence axis is highly reliable ($R = 0.94/0.96$) but a single vector captures only $\cos^2 = 0.20/0.25$ of any scenario’s own neg→pos contrast. Adding the markedness axis (the shared plane) raises the captured fraction to 0.21/0.26, which is $\approx 94\%$ of the per-scenario two-dimensional ceiling, yet the per-scenario planes are themselves spread over many dimensions (participation ratio 26.5/20.0) and the residual tilt is intent-structured (within-intent contrast cosine exceeds between-intent by 0.09/0.17). The supporting figures are Figures 15 to 17.




Figure 16: No universal steering direction (mean pooling); same panels as Figure 15. Capture plateaus near 0.25.




Figure 17: Subspace spectrum of the per-scenario centred triples. *Left*: last token (participation ratio 26.5). *Right*: mean pooling (20.0). The politeness variation is spread across many dimensions, not confined to a low-rank shared subspace.